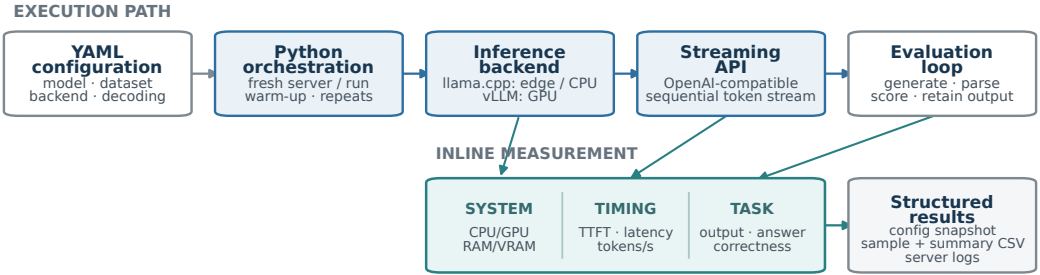
EXECUTION PATH
YAML configuration
model · dataset
backend · decoding
Python orchestration
fresh server / run
warm-up · repeats
Inference backend
llama.cpp: edge / CPU
vLLM: GPU
Streaming API
OpenAI-compatible
sequential token stream
Evaluation loop
generate · parse
score · retain output
INLINE MEASUREMENT
SYSTEM
CPU/GPU
RAM/VRAM
TIMING
TTFT · latency
tokens/s
TASK
output · answer
correctness
Structured results
config snapshot
sample + summary CSV
server logs